

# Stat 432 Homework 12

Giang Le

Assigned: Nov 15, 2021; Due: 11:59 PM CT, Nov 30, 2021

## Contents

Question 1: K-Means . . . . .	1
Question 2: Hierarchical Clustering . . . . .	7

## Question 1: K-Means

In our lecture, there is an example of clustering pixels in an image. For this question, we will replicate that procedure with your favorite image. To complete this question, perform the following steps:

- [10 Points] Pick your favorite image for this question. Plot the original image. Note that for computational concern, you probably want to avoid an extremely large image file.

```
## First I load my favorite image
install.packages("jpeg", repos='http://cran.us.r-project.org')

##
## The downloaded binary packages are in
## /var/folders/9c/3_mgdyf12z7dvv8rt4d60nt80000gn/T//RtmpnWsK8p/downloaded_packages

library(jpeg)
img<-readJPEG("/Users/gianghale/Desktop/fall-2021/stat-432/image.jpeg")

# generate a blank image
par(mar=rep(0.2, 4))

# Plot the image using rasterImage
plot(c(0, 100), c(0, 100), xaxt = 'n', yaxt = 'n',
      bty = 'n', pch = '', ylab = '', xlab = '')
rasterImage(img, 0, 0, 100, 100)
```



```
dim(img)
```

```
## [1] 500 400 3
```

- [10 Points] Report the following information of your data:
  - Dimension of the original image
  - Dimension of the data once you transform the image to a version that you could apply k-means
  - Total variations of your data

According to the code output below, the height of the original image is 500 pixels and the width of the original image is 400 pixels. The dimension of the transformed data is 200000 x3. The variation of r values from the mean is 80480.67 The variation of g values from the mean is 66516.56 The variation of b values from the mean is 37405.99

```
# Checking the dimension of the image
```

```
dim(img)
```

```
## [1] 500 400 3
```

```
# this apply function applies vectorization to each layer (r/g/b) of the image.
```

```
img_expand = apply(img, 3, c)
```

```
# and now we have the desired data matrix
```

```
dim(img_expand)
```

```
## [1] 200000 3
```

```
# Total variation of the data
```

```
means <- colMeans(img_expand)
```

```
means
```

```
## [1] 0.6780681 0.6242481 0.5160849
```

```
img_new <- setNames(as.data.frame((img_expand - means)^2), c("r", "g", "b"))
variation_r <- sum((img_new$r - means[1])^2)
variation_g <- sum((img_new$g - means[2])^2)
variation_b <- sum((img_new$b - means[3])^2)
```

```
variation_r
```

```
## [1] 80480.67
```

```
variation_g
```

```
## [1] 66516.56
```

```
variation_b
```

```
## [1] 37405.99
```

- [25 Points] Apply  $k$ -means to your data. Choose three unique  $k$  values to report the following result:
  - What is the within-cluster variance?
  - What are the cluster means?
  - Plot the image with each pixel replaced by its corresponding cluster mean

Use  $k = 5$  (answers about variance and cluster means can be found in the code chunk.)

```
set.seed(432)
# choose value = 5
kmeanfit <- kmeans(img_expand, 5)
```

```
# the within cluster variance
print("within cluster variance")
```

```
## [1] "within cluster variance"
```

```
kmeanfit$withinss
```

```
## [1] 855.7218 790.5235 3826.6861 569.3825 1308.8383
```

```
# The cluster means
print("cluster means")
```

```
## [1] "cluster means"
```

```
kmeanfit$centers
```

```
##      [,1]      [,2]      [,3]
## 1 0.7824658 0.6590773 0.53826643
## 2 0.4145895 0.2649142 0.08411499
## 3 0.7211284 0.7974984 0.78790124
## 4 0.5187149 0.4048987 0.30913240
## 5 0.8127339 0.5180558 0.08675361
```


```
# to produce the new graph, we simply replicate the cluster mean
# for all observations in the same cluster
```

```
new_img_expand = kmeanfit$centers[kmeanfit$cluster, ]
new_img = img
new_img[, , 1] = matrix(new_img_expand[,1], 500, 400)
new_img[, , 2] = matrix(new_img_expand[,2], 500, 400)
new_img[, , 3] = matrix(new_img_expand[,3], 500, 400)
```

```
# plot the new image
```

```
plot(c(0, 500), c(0, 400), xaxt = 'n', yaxt = 'n', bty = 'n',
     pch = '', ylab = '', xlab = '')

rasterImage(new_img, 0, 0, 500, 400)
```



Use  $k = 10$  (answers about variance and cluster means can be found in the code chunk.)

```
set.seed(432)
# choose value = 10
kmeanfit <- kmeans(img_expand, 10)

# the within cluster variance
print("within cluster variance")

## [1] "within cluster variance"

kmeanfit$withinss

## [1] 203.4624 364.6383 260.0683 260.8500 147.1165 361.7269 211.0988 252.1232
## [9] 616.8001 358.8716

# The cluster means
print("cluster means")

## [1] "cluster means"

kmeanfit$centers

##          [,1]      [,2]      [,3]
## 1  0.6723696 0.5545764 0.43943983
## 2  0.8728248 0.3327254 0.14873877
## 3  0.7954794 0.6774724 0.56754412
## 4  0.4985244 0.3726896 0.27669844
## 5  0.5460196 0.4865138 0.01958030
## 6  0.8442459 0.8968626 0.89388283
## 7  0.8958480 0.8032947 0.69554223
## 8  0.8339975 0.6654545 0.05139966
## 9  0.5638425 0.7418970 0.77636392
## 10 0.3638726 0.2088434 0.09038110

# to produce the new graph, we simply replicate the cluster mean
# for all observations in the same cluster
```

```

new_img_expand = kmeanfit$centers[kmeanfit$cluster, ]
new_img = img
new_img[, , 1] = matrix(new_img_expand[,1], 500, 400)
new_img[, , 2] = matrix(new_img_expand[,2], 500, 400)
new_img[, , 3] = matrix(new_img_expand[,3], 500, 400)

# plot the new image
plot(c(0, 500), c(0, 400), xaxt = 'n', yaxt = 'n', bty = 'n',
     pch = '', ylab = '', xlab = '')

rasterImage(new_img, 0, 0, 500, 400)

```



Use  $k = 25$  (answers about variance and cluster means can be found in the code chunk.)

```

set.seed(432)
# choose value = 25
kmeanfit <- kmeans(img_expand, 25)

# the within cluster variance
print("within cluster variance")

## [1] "within cluster variance"

kmeanfit$withinss

## [1] 72.80888 86.25547 73.86611 64.97486 64.66541 16.10999 61.91153
## [8] 99.97256 14.07507 105.44710 55.57529 13.37452 31.40797 65.90808
## [15] 12.25570 58.38859 56.74061 13.16350 71.07769 13.41949 59.23446
## [22] 32.45451 62.84779 16.08052 13.36083

# The cluster means
print("cluster means")

## [1] "cluster means"

kmeanfit$centers

##           [,1]      [,2]      [,3]
## 1  0.6990565 0.5937359 0.48483595
## 2  0.7539400 0.3701013 0.08334350
## 3  0.9089017 0.4770738 0.26509043

```

```
## 4  0.6169612 0.4944280 0.38240254
## 5  0.6529496 0.6112382 0.02139511
## 6  0.6099235 0.7670169 0.79892322
## 7  0.9131322 0.8038287 0.67915298
## 8  0.8864573 0.6908820 0.03292191
## 9  0.5307237 0.7201923 0.75853996
## 10 0.2906379 0.2065425 0.09128589
## 11 0.9305422 0.6773217 0.43991538
## 12 0.4441356 0.6735545 0.71618544
## 13 0.9649314 0.8888167 0.77078726
## 14 0.5758398 0.1244568 0.07683820
## 15 0.6778480 0.8118765 0.83745068
## 16 0.5115160 0.3953683 0.30573056
## 17 0.7714106 0.6820165 0.58967393
## 18 0.7573674 0.8589901 0.87536184
## 19 0.4579881 0.3301066 0.22933152
## 20 0.8668835 0.9127718 0.92032931
## 21 0.4573945 0.4162018 0.01523977
## 22 0.8102812 0.7568263 0.70906219
## 23 0.9577024 0.2684481 0.13570504
## 24 0.3207576 0.6167025 0.65611172
## 25 0.9890208 0.9784419 0.94736203

# to produce the new graph, we simply replicate the cluster mean
# for all observations in the same cluster
new_img_expand = kmeanfit$centers[kmeanfit$cluster, ]
new_img = img
new_img[, , 1] = matrix(new_img_expand[,1], 500, 400)
new_img[, , 2] = matrix(new_img_expand[,2], 500, 400)
new_img[, , 3] = matrix(new_img_expand[,3], 500, 400)

# plot the new image
plot(c(0, 500), c(0, 400), xaxt = 'n', yaxt = 'n', bty = 'n',
     pch = '', ylab = '', xlab = '')

rasterImage(new_img, 0, 0, 500, 400)
```



## Question 2: Hierarchical Clustering

The same type of image compression approach can be done using hierarchical clustering. Using the data that you prepared for the  $k$ -means algorithm, to perform hierarchical clustering. However, instead of using the euclidean distance with `dist()` function, you need to provide the clustering algorithm a different distance matrix  $D_{n \times n}$ . The  $(i, j)$ th element in this matrix represents the distance between observations  $i$  and  $j$ , defined as

$$d(\mathbf{x}_i, \mathbf{x}_k) = \|\mathbf{x}_i - \mathbf{x}_j\|_1$$

To be able to use this matrix in the `hclust()` function, you need to convert the matrix into a `dist` object, using the `as.dist()` function. For more details, read the documentation [here](#).

```
# Preparing the data.  
# Use a smaller image in order to avoid vector memory error.  
  
img1<-readJPEG("/Users/gianghale/Desktop/fall-2021/stat-432/download.jpeg")  
  
# generate a blank image  
par(mar=rep(0.2, 4))  
  
# Choosing only 100 pixels to avoid memory error.  
  
img1 <- img1[1:100,1:100,]  
  
#Plot the image  
plot(c(0, 100), c(0, 100), xaxt = 'n', yaxt = 'n',  
      bty = 'n', pch = '', ylab = '', xlab = '')  
rasterImage(img1, 0, 0, 100, 100)
```



```
img1_expand = apply(img1, 3, c)
dim(img1_expand)
```


```
## [1] 10000      3
```

```
# Complete linkage
dim(img1_expand)
```

```
## [1] 10000      3
```

```
complete <- hclust(dist(img1_expand, method = "manhattan"), method = "complete")
plot(complete)
```

## Cluster Dendrogram



```
dist(img1_expand, method = "manhattan")
hclust (*, "complete")
```

```
# Single linkage
dim(img1_expand)
```

```
## [1] 10000      3
```

```
sg <- hclust(dist(img1_expand, method = "manhattan"), method = "single")
plot(sg)
```



## Cluster Dendrogram

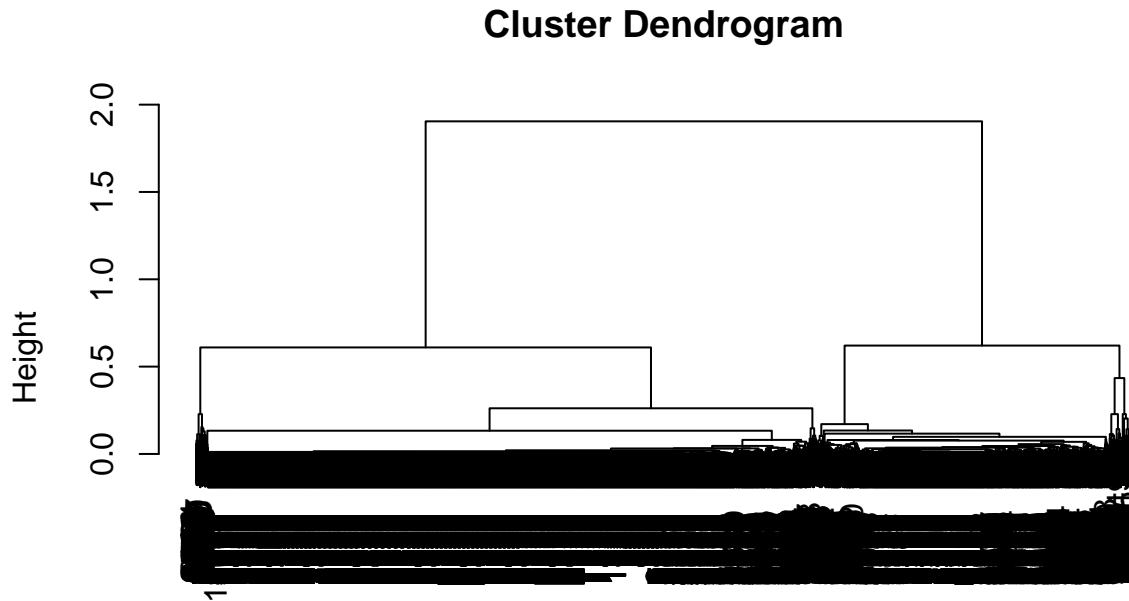


```
dist(img1_expand, method = "manhattan")  
hclust (*, "single")
```

```
# Average linkage  
dim(img1_expand)
```

```
## [1] 10000      3
```

```
ave <- hclust(dist(img1_expand, method = "manhattan"), method = "ave")  
plot(ave)
```



```
dist(img1_expand, method = "manhattan")
hclust (*, "average")
```

Once you have all the component to perform the hierarchical clustering, do the following:

- [15 Points] Try both complete, single and average linkage. Provide a plot of the dendrogram for both methods.
- [10 Points] Based on what you have, pick one final clustering result. You need to explain the rational for your choice.
- [10 Points] Based on your final choice, calculate the cluster centers using the mean of all pixels in the cluster. Then replace all pixels in each cluster with their corresponding cluster mean. This step is similar to the k-means question.
- [10 Points] Plot this new image.

Based on what you have, pick one final clustering result. You need to explain the rational for your choice. -> Looking at the plots of both the complete linkage and the average linkage, we can see that two clusters is our best choice. The reason is because if we look at the height of the split, the height before merging the two last groups is small but if we merge them, the dendrogram's height is quite large. So we can stop at two clusters. I perform clustering using k-means below with  $k = 2$  and plotted the new image below.

```
set.seed(432)
# choose value = 2
kmeanfit <- kmeans(img1_expand, 2)

# The cluster means
print("cluster means")

## [1] "cluster means"
kmeanfit$centers

##          [,1]      [,2]      [,3]
## 1 0.08591672 0.3562926 0.6255952
## 2 0.98866592 0.9924638 0.9940001
```

```

# to produce the new graph, we simply replicate the cluster mean
# for all observations in the same cluster
new_img_expand = kmeanfit$centers[kmeanfit$cluster, ]
new_img = img1
new_img[, , 1] = matrix(new_img_expand[,1], 100, 100)
new_img[, , 2] = matrix(new_img_expand[,2], 100, 100)
new_img[, , 3] = matrix(new_img_expand[,3], 100, 100)

# plot the new image
plot(c(0, 100), c(0, 100), xaxt = 'n', yaxt = 'n', bty = 'n',
     pch = '', ylab = '', xlab = '')

rasterImage(new_img, 0, 0, 100, 100)

```

